



OCaml

FROM IMPERATIVE TO FUNCTIONAL APIS

Marek Kubica @leonidasfromxiv


28. March 2015



stylefruits

- They sponsored me going to this conference
- Jan Stępień is holding a talk in the currying hall at 11:00
- ... and a workshop at 14:00!
- Thanks, stylefruits!

- Marek Kubica
- Student at the TUM
- I do free software
- Dabbled in just about every language ever



- Data compression support in OCaml is... kinda meh
- Using OCaml foreign function interface to talk to libarchive
- Make it a bachelor thesis!
- Thought I might as well create a better API
- What *IS* a better API anyway?

- Data compression support in OCaml is... kinda meh
- Using OCaml foreign function interface to talk to libarchive
- Make it a bachelor thesis!
- Thought I might as well create a better API
- What *IS* a better API anyway?

My goal for today

Show you that advanced static type features are not (only) academic.

Let's check how C handles look like. How does libarchive handle this?

```
__LA_DECT struct archive* archive_read_new( void );  
      C return type      Function name      Argument type  
__LA_DECT struct archive* archive_write_new( void );  
      C return type      Function name      Argument type
```

- Opaque pointer to some struct
- Write handles and read handles have the same type

So, what can we do with these handles?

So, what can we do with these handles?

- Create them
- Open them
- Configure them
- Read from them
- Write to them
- Close them

So, what can we do with these handles?

- Create them
- Open them
- Configure them
- Read from them
- Write to them
- Close them

Cool.

So, what can we do with these handles?

- Create them
- Open them
- Configure them
- Read from them
- Write to them
- Close them

Cool. But what if we screw up?

```
zsh: segmentation fault (core dumped) ./errors
```

```

*** Error in './errors': double free or corruption (fasttop): 0x000000000077a010 ***
===== Backtrace: =====
/usr/lib/libc.so.6(+0x788ae)[0x7fa0c97cd8ae]
/usr/lib/libc.so.6(+0x79587)[0x7fa0c97ce587]
./errors[0x40057b]
/usr/lib/libc.so.6(__libc_start_main+0xf5)[0x7fa0c9776a15]
./errors[0x400479]
===== Memory map: =====
00400000-00401000 r-xp 00000000 fe:01 21244012 /lambdacon/errors
00600000-00601000 rw-p 00000000 fe:01 21244012 /lambdacon/errors
0077a000-0079b000 rw-p 00000000 00:00 0 [heap]
7fa0c953f000-7fa0c9554000 r-xp 00000000 fe:00 4213606 /usr/lib/libgcc_s.so.1
7fa0c9554000-7fa0c9754000 ---p 00015000 fe:00 4213606 /usr/lib/libgcc_s.so.1
7fa0c9754000-7fa0c9755000 rw-p 00015000 fe:00 4213606 /usr/lib/libgcc_s.so.1
7fa0c9755000-7fa0c98f8000 r-xp 00000000 fe:00 4202529 /usr/lib/libc-2.17.so
7fa0c98f8000-7fa0c9af8000 ---p 001a3000 fe:00 4202529 /usr/lib/libc-2.17.so
7fa0c9af8000-7fa0c9afc000 r---p 001a3000 fe:00 4202529 /usr/lib/libc-2.17.so
7fa0c9afc000-7fa0c9afe000 rw-p 001a7000 fe:00 4202529 /usr/lib/libc-2.17.so
7fa0c9afe000-7fa0c9b02000 rw-p 00000000 00:00 0
7fa0c9b02000-7fa0c9b23000 r-xp 00000000 fe:00 4203728 /usr/lib/ld-2.17.so
7fa0c9b23000-7fa0c9cfd000 rw-p 00000000 00:00 0
7fa0c9cfd000-7fa0c9d23000 rw-p 00000000 00:00 0
7fa0c9d23000-7fa0c9d24000 r---p 00021000 fe:00 4203728 /usr/lib/ld-2.17.so
7fa0c9d24000-7fa0c9d25000 rw-p 00022000 fe:00 4203728 /usr/lib/ld-2.17.so
7fa0c9d25000-7fa0c9d26000 rw-p 00000000 00:00 0
7fff44461000-7fff44482000 rw-p 00000000 00:00 0 [stack]
7fff445fe000-7fff44600000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 r-xp 00000000 00:00 0 [vsyscall]
zsh: abort (core dumped) ./errors foo

```

What actually happens: libarchive returns **ARCHIVE_FATAL**.

What actually happens: libarchive returns **ARCHIVE_FATAL**.

Unless you trigger a bug in libarchive.

What actually happens: libarchive returns **ARCHIVE_FATAL**.

Unless you trigger a bug in libarchive. **Then it segfaults.**

Lots of things can go wrong

- Reading from handle that is not open *
- Writing to a read handle
- Not setting the options correctly (compression formats)

Lots of things can go wrong

- Reading from handle that is not open *
- Writing to a read handle
- Not setting the options correctly (compression formats)

* this actually happened

Not to gripe on libarchive...

Public Service Announcement

libarchive is a rather well designed library. Mostly idiomatic C, so don't think this is a deliberately bad example. It is how things *are* in C land.

- This is an OK API for C.
- Fragile APIs are common in C.

Not to gripe on libarchive...

Public Service Announcement

libarchive is a rather well designed library. Mostly idiomatic C, so don't think this is a deliberately bad example. It is how things *are* in C land.

- This is an OK API for C.
- Fragile APIs are common in C.

But can we do better?

Yes

Yes

(obviously)

How to prevent writing to read handles and reading from write handles?

```
external read_new: unit -> archive = "ost_read_new"  
external write_new: unit -> archive = "ost_write_new"
```

FIX UP THE HANDLE TYPES

How to prevent writing to read handles and reading from write handles?

```
external read_new: unit -> archive = "ost_read_new"  
external write_new: unit -> archive = "ost_write_new"
```

Yup, create different handle types.

```
type r = archive  
type w = (archive * write_buffer_ptr * written_ptr)
```

So now we have distinct types to represent handles.




Type safety improved
Writing read handles disallowed

But of course you aren't attending the talk for this **trivial** epiphany.
We can do this easily in C as well! Let's do better.

OUR HANDLES HAVE STATES!

The handles always traverse some fixed states:



Couldn't we **encode the state** in the type somehow?

We can add state. OCaml has parametrized types*:

```
# [];;
- : 'a list = []
# [1];;
- : int list = [1]
# type 'a read_handle = ReadHandle of 'a;;
type 'a read_handle = ReadHandle of 'a
```

* if you haven't seen them, think of them kinda like generics

So, now we can **parametrize types with other types**.

We could create our own state types:

```
type state = New | Configured | Opened | Closed
```

So, now we can **parametrize types with other types**.

We could create our own state types:

```
type state = New | Configured | Opened | Closed
```

But these **can't be extended** if someone wants to add a new state.

So, now we can **parametrize types with other types**.

We could create our own state types:

```
type state = New | Configured | Opened | Closed
```

But these **can't be extended** if someone wants to add a new state.

Plus, we're lazy. Let's use **open union types** aka polymorphic variants:

```
[ 'New ] [ 'Configured ] [ 'Opened ] [ 'Closed ]
```

Great, so now we can create functions that *require* handles of the correct state.

e.g. a **read** function that only works on [**'Opened'**]
read_handle.

Great, so now we can create functions that *require* handles of the correct state.

e.g. a **read** function that only works on ['Opened'
read_handle.

Foiled again!

Great, so now we can create functions that *require* handles of the correct state.

e.g. a **read** function that only works on `['Opened]`
read_handle.

Foiled again!

The OCaml compiler is too smart, it knows that `['Opened]`
read_handle is the same type as `['New] read_handle`.
Therefore every function which takes the `['Opened]` handle
accepts every other type of handle as well.

We'd need to **hide** the actual `read_handle` type from the compiler.

We'd need to **hide** the actual `read_handle` type from the compiler.

Boy oh boy, we can!

We'd need to **hide** the actual `read_handle` type from the compiler.

Boy oh boy, we can!

We create a module and only say:

```
module Handle : sig  
  type 'a r  
  (* our signatures *)  
  val new : unit -> ['New] r  
end = struct  
  type 'a r = read_handle  
  (* our functions *)  
  external new : unit -> ['Open] r = "ost_read_new"  
end
```



Type safety improved
Made API misuse a type error

- ✓ Writing read handles disallowed
- ✓ Using the proper handle in an incorrect way disallowed

And now for something
completely different!

HAVE YOU EVER SEEN THIS?

Ever used Python?

```
Traceback (most recent call last):  
File "<stdin>", line 1, in <module>  
AttributeError: 'NoneType' object has no attribute 'foo'
```

HAVE YOU EVER SEEN THIS?

Ever used Python?

```
Traceback (most recent call last):  
File "<stdin>", line 1, in <module>  
AttributeError: 'NoneType' object has no attribute 'foo'
```

Ever touched Java?

```
Exception in thread "main" java.lang.NullPointerException  
at NPE.main(NPE.java:8)
```

HAVE YOU EVER SEEN THIS?

Ever used Python?

```
Traceback (most recent call last):  
File "<stdin>", line 1, in <module>  
AttributeError: 'NoneType' object has no attribute 'foo'
```

Ever touched Java?

```
Exception in thread "main" java.lang.NullPointerException  
at NPE.main(NPE.java:8)
```

Ever seen C?

```
zsh: segmentation fault (core dumped) ./errors
```


You know whose fault it is!

- `null`
- `None`
- `NULL`
- `nil`

You know whose fault it is!

- `null`
- `None`
- `NULL`
- `nil`

Everytime you return `NULL` as a placeholder value, ~~`$DEITY` kills a kitten~~ you have to check whether you weren't handed `NULL` in return.

Let's kill the ~~Batman~~ Null
pointer!

- Common solution
- Ubiquitous (Java, Python, C++, Ruby, whathaveyou)
- Easy to understand
- OCaml does have exceptions

- Common solution
- Ubiquitous (Java, Python, C++, Ruby, whathaveyou)
- Easy to understand
- OCaml does have exceptions
- **Not typesafe**, unless you consider checked exceptions
- Boring!

Observation: when we return NULL, we either return something meaningful or a marker that there was nothing to return.

We might even say:

```
type 'a option = Some of 'a | None
```

ATTEMPT TWO: OPTION TYPES

Observation: when we return NULL, we either return something meaningful or a marker that there was nothing to return.

We might even say:

```
type 'a option = Some of 'a | None
```

Therefore, everytime a function returns `'a option` we have to pattern match:

```
let optional x = Some x  
match optional 42 with  
  | Some x -> x  
  | None -> 0
```

ATTEMPT TWO: OPTION TYPES

Observation: when we return NULL, we either return something meaningful or a marker that there was nothing to return.

We might even say:

```
type 'a option = Some of 'a | None
```

Therefore, everytime a function returns `'a option` we have to pattern match:

```
let optional x = Some x  
match optional 42 with  
  | Some x -> x  
  | None -> 0
```

If we forget:



Type safety improved
Missed null check is type error

✓ No more Null pointer failures on runtime!

Everything is fun and games until you need to specify a **reason** for failure.

What if we could add an error message?

```
type ('a, 'b) err = Success of 'a | Failure of 'b
```

Done!

But pattern matching on every function call sucks because it is tedious! Just look at this mess:

```
match firstfn 42 with
| Success (x) -> (match secondfn x with
| Success (y) -> (match thirdfn y with
| Success (z) -> z
| Failure (f3) -> "Failure at thirdfn")
| Failure (f2) -> "Failure at secondfn")
| Failure (f1) -> "Failure at fristfn"
```

Right. Maybe we can simplify...

In Haskell, **option** is called "Maybe monad" and **error** is called "Error monad".

But pattern matching on every function call sucks because it is tedious! Just look at this mess:

```
match firstfn 42 with
| Success (x) -> (match secondfn x with
| Success (y) -> (match thirdfn y with
| Success (z) -> z
| Failure (f3) -> "Failure at thirdfn")
| Failure (f2) -> "Failure at secondfn")
| Failure (f1) -> "Failure at fristfn"
```

Right. Maybe we can simplify...

In Haskell, **option** is called “Maybe monad” and **error** is called “Error monad”.

BAM, SCARY MONADS!

Haskell features an operator called **bind** aka **»=** to chain operations on monads.

```
val bind: 'a ErrorMonad.t ->  
  ('a -> 'b ErrorMonad.t) ->  
  'b ErrorMonad.t
```

bind takes an error monad wrapping type **'a**, and a function which takes **'a** and returns an error monad wrapping **'b** and returns that value.

Basically an unwrapper function.

For the error monad, it looks like this:

```
let bind m f = match m with  
  | Success(x) -> f x  
  | Failure(f)  -> Failure(f)
```

We can use it like this:

```
match (bind (bind (firstfn 42) secondfn) thirdfn)  
with  
  | Success (x) -> x  
  | Failure (_) -> "Failure in chain"
```

The code got a lot easier!

OCaml allows **custom operators** as long as they follow naming rules.

```
let (»=) = bind
```

Using it is **easy**:

```
match (firstfn 42) »= secondfn »= thirdfn with  
  | Success (x) -> x  
  | Failure (_) -> "Failure in chain"
```



Type safety improved Statically typed error handling

- ✓ No more Null pointer failures on runtime!
- ✓ Easy and convenient to get reason of failure

Sometimes, errors will happen

```
let divide a b =  
  match b with  
  | 0 -> Failure "division"  
  | b -> Success (a / b)  
  
let handle_user_input () =  
  match divide 42 (read_int ()) with  
  | Success res -> Printf.sprintf "Got %d" res  
  | Failure "division_by_zero" -> "Divided by zero"
```

Can you spot the error?

Sometimes, errors will happen

```
let divide a b =  
  match b with  
  | 0 -> Failure "division"  
  | b -> Success (a / b)  
  
let handle_user_input () =  
  match divide 42 (read_int ()) with  
  | Success res -> Printf.sprintf "Got %d" res  
  | Failure "division_by_zero" -> "Divided by zero"
```

Can you spot the error?

```
divide : int -> int -> (int, string) err
```

We could define a type constructor for each error case:

```
type division_error = Division_by_zero | Overflow
```

This **works** in this case, but what if we want to **reuse** constructors?

```
type multiplication_error = Overflow
```

Does not compile. Each constructor can only be of *one* type.

Polymorphic variant constructors can be composed into types:

```
type division_error = [  
  | 'Division_by_zero  
  | 'Overflow  
]
```

Works like sets. OCaml does it automatically if functions return variants.

Let's fix the program

```
let divide a b =  
  match b with  
  | 0 -> Failure 'Division'  
  | b -> Success (a / b)
```

```
let handle_user_input () =  
  match divide 42 (read_int ()) with  
  | Success res -> Printf.sprintf "Got %d" res  
  | Failure 'Division -> "Divided by zero"
```

```
divide: int -> int -> (int, [> 'Division ]) err  
                                our error variant
```



Achievement unlocked
Possible errors reflected in type system

- ✓ Possible errors can be seen in signatures
- ✓ Type system can warn when errors not handled

There are many more tricks on how you can use the type system, to make illegal state unrepresentable, e.g. Generalized Algebraic Data Types (GADTs).

But take care: the API might turn out to be *too complicated*. Please, use common sense*.

There are many more tricks on how you can use the type system, to make illegal state unrepresentable, e.g. Generalized Algebraic Data Types (GADTs).

But take care: the API might turn out to be *too complicated*. Please, use common sense*.

* if not applicable, emulate idioms from good APIs in your preferred programming language ☺

Marek Kubica

Check out my playthings:

- Leonidas-from-XIV on GitHub
- @leonidasfromxiv
- <https://xivilization.net/>

